

Video Game Analysis

January 23, 2024

```
[117]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
from scipy.stats import probplot
import statsmodels.api as sm
import statsmodels.formula.api as smf
from sklearn.preprocessing import LabelEncoder
from sklearn import metrics
from sklearn import linear_model
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.model_selection import cross_val_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.stats.api as sms
import statsmodels.stats.outliers_influence as oi
import scipy.stats as stats
from boruta import BorutaPy
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

1 Import data files

All the data belongs to Kaggle

```
[118]: video_games_sales = pd.read_csv('video_games_sales.csv')

video_games_sales
```

```
[118]:      rank      name platform \
0      1      Wii Sports      Wii
```

1	2	Super Mario Bros.	NES
2	3	Mario Kart Wii	Wii
3	4	Wii Sports Resort	Wii
4	5	Pokemon Red/Pokemon Blue	GB
...
16593	16596	Woody Woodpecker in Crazy Castle 5	GBA
16594	16597	Men in Black II: Alien Escape	GC
16595	16598	SCORE International Baja 1000: The Official Game	PS2
16596	16599	Know How 2	DS
16597	16600	Spirits & Spells	GBA

	year	genre	publisher	na_sales	eu_sales	jp_sales	\
0	2006.0	Sports	Nintendo	41.49	29.02	3.77	
1	1985.0	Platform	Nintendo	29.08	3.58	6.81	
2	2008.0	Racing	Nintendo	15.85	12.88	3.79	
3	2009.0	Sports	Nintendo	15.75	11.01	3.28	
4	1996.0	Role-Playing	Nintendo	11.27	8.89	10.22	
...	
16593	2002.0	Platform	Kemco	0.01	0.00	0.00	
16594	2003.0	Shooter	Infogrames	0.01	0.00	0.00	
16595	2008.0	Racing	Activision	0.00	0.00	0.00	
16596	2010.0	Puzzle	7G//AMES	0.00	0.01	0.00	
16597	2003.0	Platform	Wanadoo	0.01	0.00	0.00	

	other_sales	global_sales
0	8.46	82.74
1	0.77	40.24
2	3.31	35.82
3	2.96	33.00
4	1.00	31.37
...
16593	0.00	0.01
16594	0.00	0.01
16595	0.00	0.01
16596	0.00	0.01
16597	0.00	0.01

[16598 rows x 11 columns]

1.0.1 Video Game Sales

- rank = Sales rank
- name = Name of video game
- platform = Video Game Platform
- year = When game was released
- genre = Game Genre
- publisher = Video Game Publisher
- na_sales = North America Sales, Millions

- eu_sales = European Sales, Millions
- jp_sales = Japan Sales, Millions
- other_sales = Sales in other areas, millions

```
[119]: data = video_games_sales.copy()
```

2 Clean Data

Check for missing values

- If missing values found, impute or drop them based off data
- In this case, we will drop the missing values because you can not impute the year or publisher

```
[120]: data.isnull().sum()
```

```
[120]: rank          0
      name          0
      platform      0
      year         271
      genre         0
      publisher     58
      na_sales      0
      eu_sales      0
      jp_sales      0
      other_sales   0
      global_sales  0
      dtype: int64
```

```
[121]: data = data.dropna()
      data.isnull().sum()
```

```
[121]: rank          0
      name          0
      platform      0
      year          0
      genre         0
      publisher     0
      na_sales      0
      eu_sales      0
      jp_sales      0
      other_sales   0
      global_sales  0
      dtype: int64
```

Check for duplicates

- If duplicates found, we drop them
- No duplicates were found

```
[122]: data.duplicated().sum()
```

```
[122]: 0
```

Fix data types

```
[123]: # Convert the columns below into categorical data
data['genre'] = data['genre'].astype('category')
data['publisher'] = data['publisher'].astype('category')
data['platform'] = data['platform'].astype('category')

# Convert float to datetime year
data['year'] = pd.to_datetime(data['year'], format='%Y', errors='coerce')
data['year'] = data['year'].dt.year
```

```
[124]: data.rename(columns={'na_sales': 'North America Sales',
                           'eu_sales': 'European Sales',
                           'jp_sales': 'Japan Sales',
                           'other_sales': 'Other Sales',
                           'global_sales': 'Global Sales'}, inplace=True)
```

3 Exploratory Data Analysis & Ad-hoc Analysis

Check data summary

```
[125]: data.describe()
```

```
[125]:
```

	rank	year	North America Sales	European Sales	\
count	16291.000000	16291.000000	16291.000000	16291.000000	
mean	8290.190228	2006.405561	0.265647	0.147731	
std	4792.654450	5.832412	0.822432	0.509303	
min	1.000000	1980.000000	0.000000	0.000000	
25%	4132.500000	2003.000000	0.000000	0.000000	
50%	8292.000000	2007.000000	0.080000	0.020000	
75%	12439.500000	2010.000000	0.240000	0.110000	
max	16600.000000	2020.000000	41.490000	29.020000	

	Japan Sales	Other Sales	Global Sales
count	16291.000000	16291.000000	16291.000000
mean	0.078833	0.048426	0.540910
std	0.311879	0.190083	1.567345
min	0.000000	0.000000	0.010000
25%	0.000000	0.000000	0.060000
50%	0.000000	0.010000	0.170000
75%	0.040000	0.040000	0.480000
max	10.220000	10.570000	82.740000

```
[126]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 16291 entries, 0 to 16597
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rank                   16291 non-null  int64
1   name                   16291 non-null  object
2   platform               16291 non-null  category
3   year                   16291 non-null  int32
4   genre                  16291 non-null  category
5   publisher              16291 non-null  category
6   North America Sales    16291 non-null  float64
7   European Sales         16291 non-null  float64
8   Japan Sales            16291 non-null  float64
9   Other Sales            16291 non-null  float64
10  Global Sales           16291 non-null  float64
dtypes: category(3), float64(5), int32(1), int64(1), object(1)
memory usage: 1.1+ MB
```

```
[127]: data.shape
```

```
[127]: (16291, 11)
```

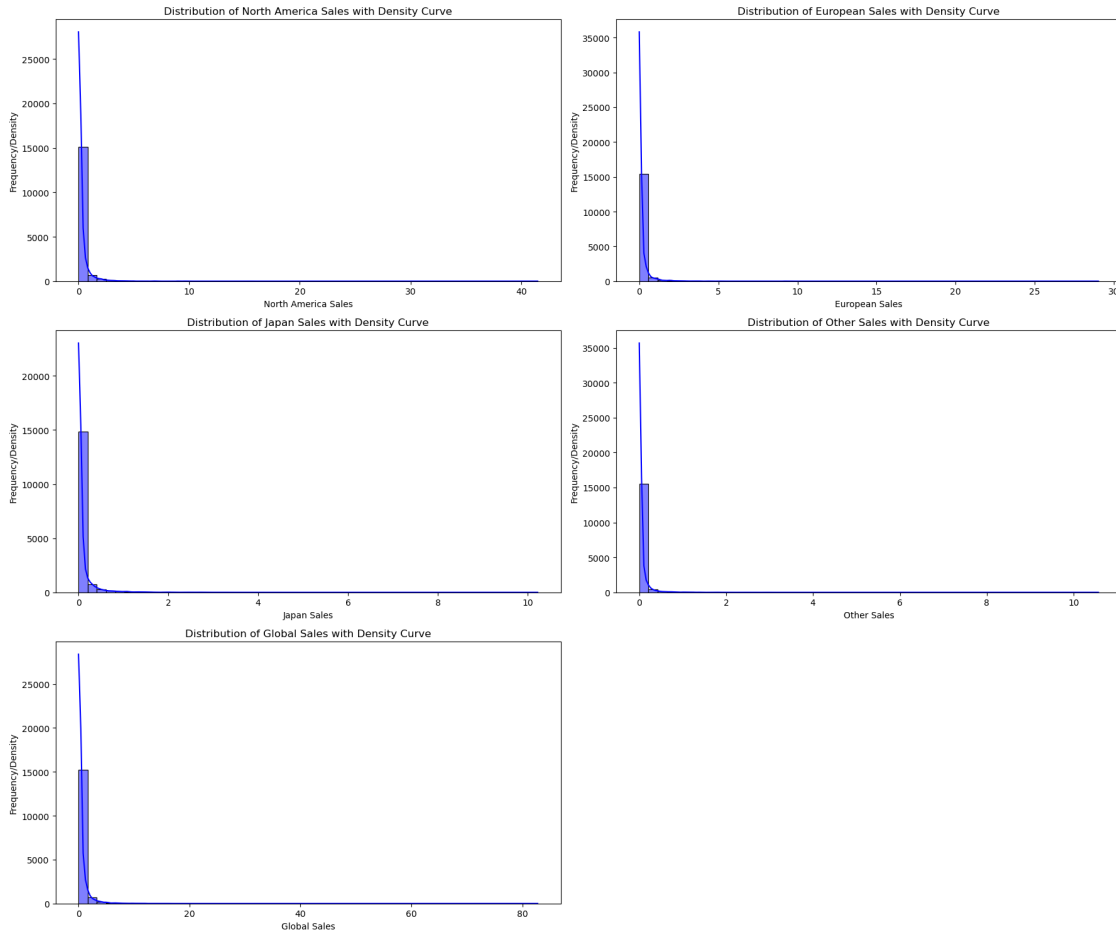
Histogram

```
[128]: columns = ['North America Sales', 'European Sales', 'Japan Sales', 'Other_
↳Sales', 'Global Sales']
```

```
[129]: plt.figure(figsize=(18, 15))

for i, column in enumerate(columns, 1):
    plt.subplot(3, 2, i)
    sns.histplot(data[column], bins=50, kde=True, color='blue',
↳edgecolor='black')
    plt.title(f'Distribution of {column} with Density Curve')
    plt.xlabel(column)
    plt.ylabel('Frequency/Density')

plt.tight_layout()
plt.show()
```




Boxplot

```
[130]: plt.figure(figsize=(18, 10))

sns.boxplot(data=data[columns], width=0.5, palette='pastel')

plt.title('Box Plot of Sales Data by Region')
plt.xlabel('Region')
plt.ylabel('Sales')
plt.show()
```




Quantile Plots

```
[131]: plt.figure(figsize=(18, 15))

for i, column in enumerate(columns, 1):
    plt.subplot(3, 2, i)
    probplot(data[column], plot=plt)
    plt.title(f'QQ Plot for {column}')


plt.tight_layout()
plt.show()
```



Correlation Matrix

```
[132]: correlation_matrix = data[['North America Sales', 'European Sales', 'Japan_
↳Sales', 'Other Sales', 'Global Sales']].corr()

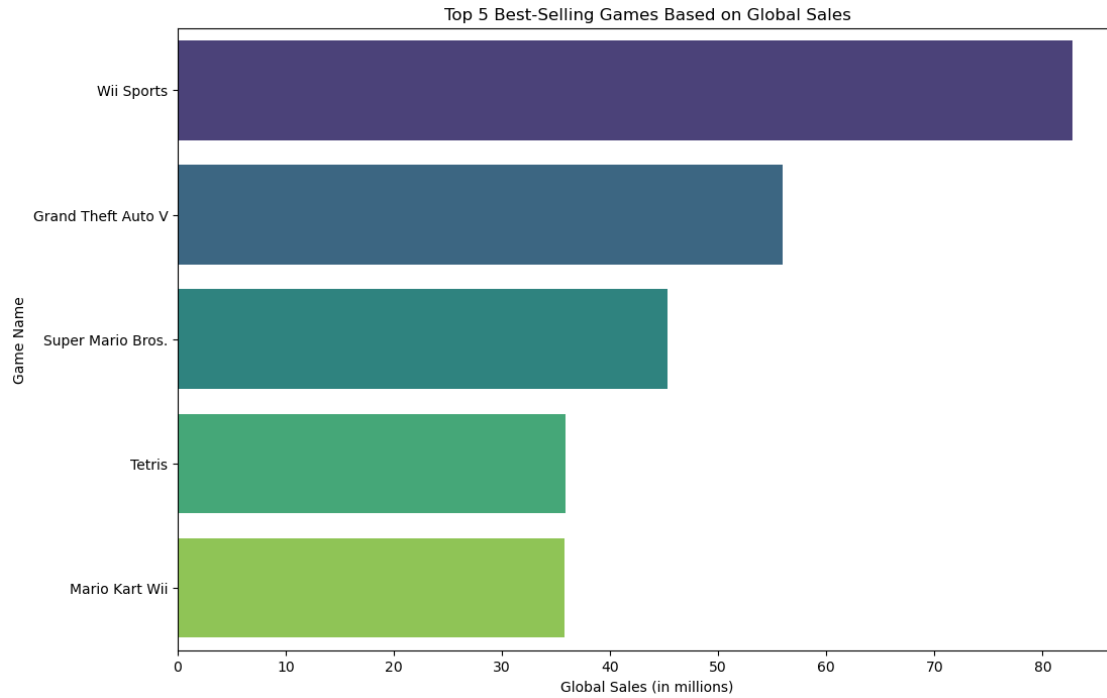
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```

Top 5 Best Seller Games

```
[133]: top_sellers_combined = data.groupby('name')['Global Sales'].sum().nlargest(5)

plt.figure(figsize=(12, 8))
sns.barplot(x=top_sellers_combined.values, y=top_sellers_combined.index,
            palette='viridis')
plt.title('Top 5 Best-Selling Games Based on Global Sales')
plt.xlabel('Global Sales (in millions)')
plt.ylabel('Game Name')
plt.show()
```

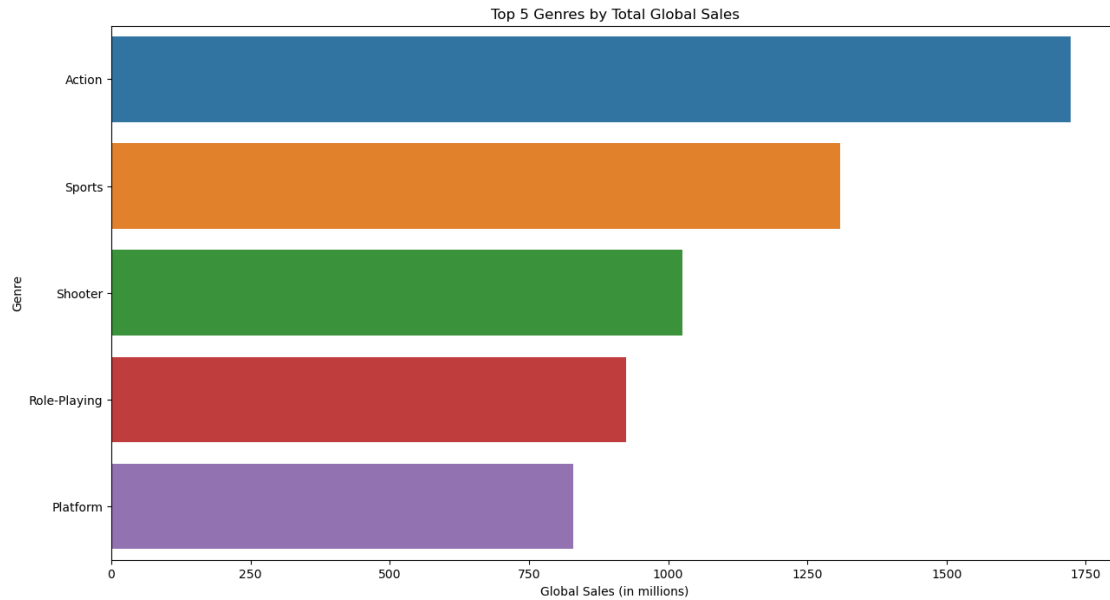


Top 5 genres

```
[134]: genre_sales = data.groupby('genre')['Global Sales'].sum().nlargest(5)

top_genres = genre_sales.index.tolist()[:5]

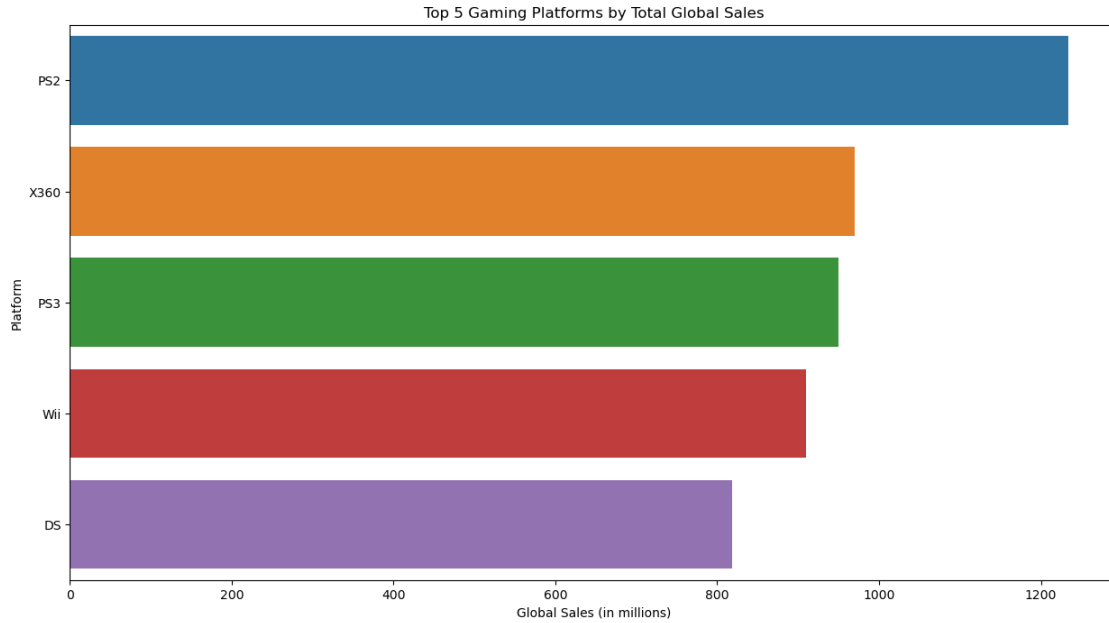
plt.figure(figsize=(15, 8))
sns.barplot(x=genre_sales.values, y=genre_sales.index, order=top_genres)
plt.title('Top 5 Genres by Total Global Sales')
plt.xlabel('Global Sales (in millions)')
plt.ylabel('Genre')
plt.show()
```



Top 5 Platforms

```
[135]: platform_sales = data.groupby('platform')['Global Sales'].sum().nlargest(5)
top_platform = platform_sales.index.tolist()[:5]

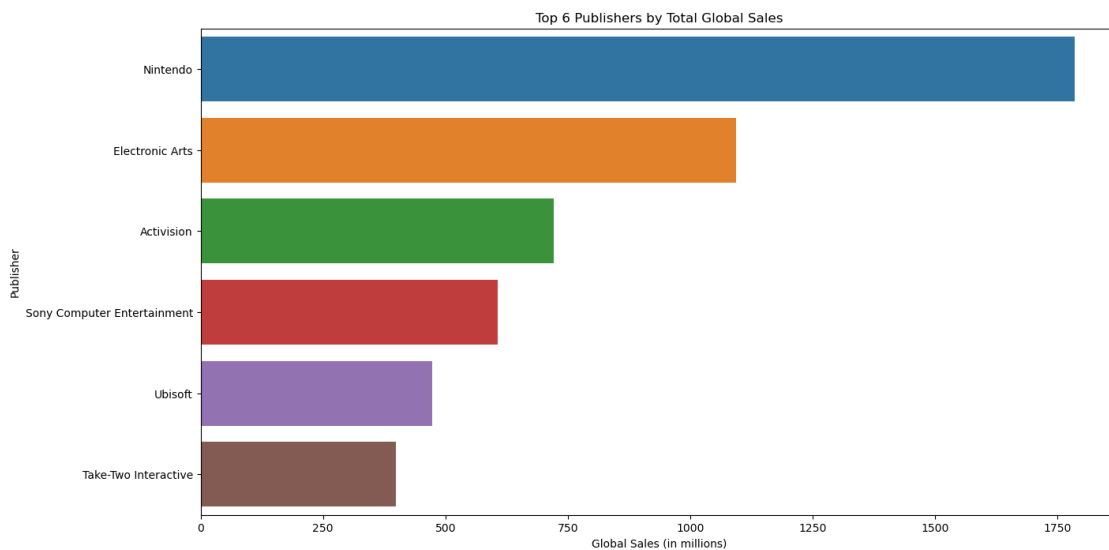
plt.figure(figsize=(15, 8))
sns.barplot(x=platform_sales.values, y=top_platform, order=top_platform)
plt.title('Top 5 Gaming Platforms by Total Global Sales')
plt.xlabel('Global Sales (in millions)')
plt.ylabel('Platform')
plt.show()
```



Top 5 publisher

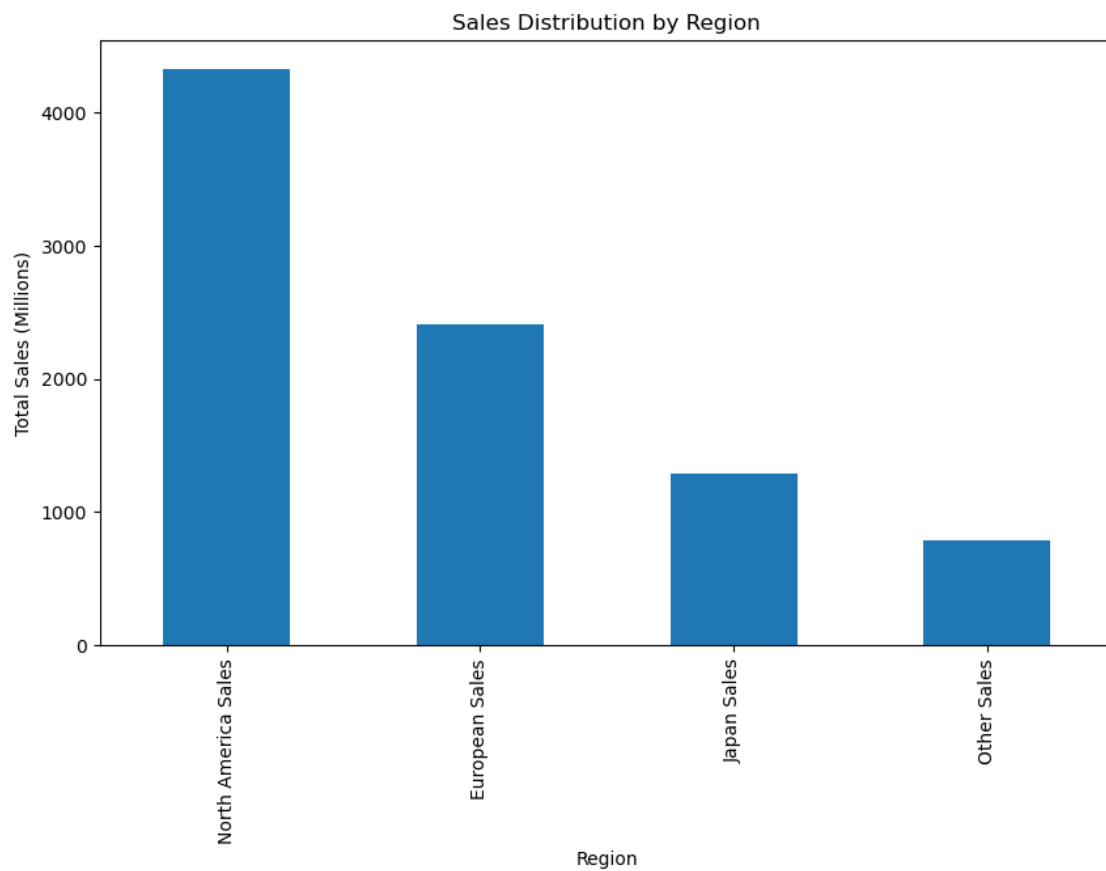
```
[136]: publisher_sales = data.groupby('publisher')['Global Sales'].sum().nlargest(6)
top_publishers = publisher_sales.index.tolist()[:6]

plt.figure(figsize=(15, 8))
sns.barplot(x=publisher_sales.values, y=top_publishers, order=top_publishers)
plt.title('Top 6 Publishers by Total Global Sales')
plt.xlabel('Global Sales (in millions)')
plt.ylabel('Publisher')
plt.show()
```



Sales Distribution by Region

```
[137]: regions = ['North America Sales', 'European Sales', 'Japan Sales', 'Other_␣  
↳Sales']  
  
plt.figure(figsize=(10, 6))  
data[regions].sum().plot(kind='bar')  
plt.title('Sales Distribution by Region')  
plt.xlabel('Region')  
plt.ylabel('Total Sales (Millions)')  
plt.show()
```



Yearly Trends

```
[138]: plt.figure(figsize=(12, 6))  
data.groupby('year')['Global Sales'].sum().plot(kind='line', marker='o', ␣  
↳color='purple')  
plt.title('Global Sales Trend 1997-2016')
```

```
plt.xlabel('Year')
plt.ylabel('Total Global Sales (Millions)')
plt.show()
```




Genre-Platform Combinations

```
[139]: genre_platform_combinations = data.pivot_table(index='genre',
    ↪ columns='platform', values='Global Sales', aggfunc='sum')

plt.figure(figsize=(12, 8))
sns.heatmap(genre_platform_combinations, cmap='viridis', fmt='.1f', linewidths=.
    ↪ 5)

plt.title('Genre-Platform Combinations - Global Sales')
plt.xlabel('Platform')
plt.ylabel('Genre')
plt.show()
```



```
[140]: genre_platform_combinations = data.pivot_table(index='genre',
columns='platform', values='Global Sales', aggfunc='sum')

top_combinations = genre_platform_combinations.stack().idxmax()

print("Most Successful Genre-Platform Combinations:")
print(top_combinations)
```

Most Successful Genre-Platform Combinations:
('Action', 'PS3')

4 A/B Testing

Platform Impact Hypothesis

```
[141]: platform_1_name = 'PS2'
platform_2_name = 'X360'

platform_1_sales = data[data['platform'] == platform_1_name]['Global Sales']
platform_2_sales = data[data['platform'] == platform_2_name]['Global Sales']

t_stat, p_value = stats.ttest_ind(platform_1_sales, platform_2_sales)
```

```

alpha = 0.05

if p_value < alpha:
    print(f"Reject the null hypothesis. There is a significant difference in_
    ↪global sales between {platform_1_name} and {platform_2_name}.")
else:
    print(f"Fail to reject the null hypothesis. There is no significant_
    ↪difference in global sales between {platform_1_name} and {platform_2_name}.")

```

Reject the null hypothesis. There is a significant difference in global sales between PS2 and X360.

Genre Impact Hypothesis:

```

[142]: genre_1_name = 'Action'
        genre_2_name = 'Shooter'

        genre_1_sales = data[data['genre'] == genre_1_name]['Global Sales']
        genre_2_sales = data[data['genre'] == genre_2_name]['Global Sales']

        t_stat_genre, p_value_genre = stats.ttest_ind(genre_1_sales, genre_2_sales)

        alpha = 0.05

        if p_value_genre < alpha:
            print(f"Reject the null hypothesis. There is a significant difference in_
            ↪global sales between {genre_1_name} and {genre_2_name}.")
        else:
            print(f"Fail to reject the null hypothesis. There is no significant_
            ↪difference in global sales between {genre_1_name} and {genre_2_name}.")

```

Reject the null hypothesis. There is a significant difference in global sales between Action and Shooter.

Publisher Impact Hypothesis:

```

[143]: publisher_1_name = 'Take-Two Interactive'
        publisher_2_name = 'Activision'

        publisher_1_sales = data[data['publisher'] == publisher_1_name]['Global Sales']
        publisher_2_sales = data[data['publisher'] == publisher_2_name]['Global Sales']

        t_stat_publisher, p_value_publisher = stats.ttest_ind(publisher_1_sales,
        ↪publisher_2_sales)

        alpha = 0.05

```



```

if p_value_publisher < alpha:
    print(f"Reject the null hypothesis. There is a significant difference in_
    ↪ global sales between {publisher_1_name} and {publisher_2_name}.")
else:
    print(f"Fail to reject the null hypothesis. There is no significant_
    ↪ difference in global sales between {publisher_1_name} and {publisher_2_name}.
    ↪")

```

Reject the null hypothesis. There is a significant difference in global sales between Take-Two Interactive and Activision.

5 Model Building

5.0.1 Fit the model

```

[144]: label_encoder = LabelEncoder()

data['genre_encoded'] = label_encoder.fit_transform(data['genre'])
data['publisher_encoded'] = label_encoder.fit_transform(data['publisher'])
data['platform_encoded'] = label_encoder.fit_transform(data['platform'])

```

```

[145]: X = data[['genre_encoded', 'publisher_encoded', 'platform_encoded']]

y = data['Global Sales']

```

```

[146]: X = sm.add_constant(X)

model = sm.OLS(y, X).fit()

print(model.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  Global Sales    R-squared:                        0.002
Model:                            OLS         Adj. R-squared:                  0.001
Method:                        Least Squares   F-statistic:                       9.146
Date:                Tue, 23 Jan 2024         Prob (F-statistic):                4.82e-06
Time:                        16:56:42         Log-Likelihood:                   -30423.
No. Observations:                16291        AIC:                        6.085e+04
Df Residuals:                    16287        BIC:                        6.088e+04
Df Model:                          3
Covariance Type:                nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
-----

```

const	0.3581	0.037	9.606	0.000	0.285
0.431					
genre_encoded	0.0094	0.003	2.876	0.004	0.003
0.016					
publisher_encoded	0.0002	6.96e-05	2.681	0.007	5.01e-05
0.000					
platform_encoded	0.0052	0.001	3.537	0.000	0.002
0.008					

```
=====
```

Omnibus:	33706.268	Durbin-Watson:	0.050
Prob(Omnibus):	0.000	Jarque-Bera (JB):	241118304.225
Skew:	17.294	Prob(JB):	0.00
Kurtosis:	597.996	Cond. No.	1.04e+03

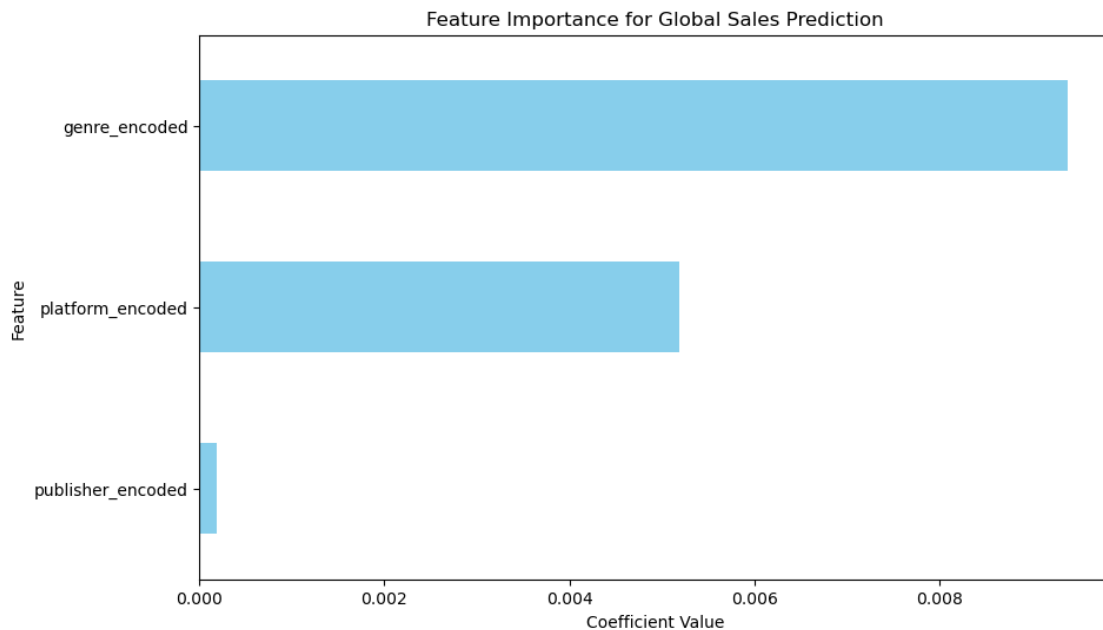
```
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.04e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[147]: feature_importance = model.params.drop("const")

plt.figure(figsize=(10, 6))
feature_importance.sort_values().plot(kind='barh', color='skyblue')
plt.title('Feature Importance for Global Sales Prediction')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')
plt.show()
```



5.0.2 Make Predictions for Global Sales

Predictions at Random

```
[148]: genre_range = data['genre_encoded'].max() - data['genre_encoded'].min()
publisher_range = data['publisher_encoded'].max() - data['publisher_encoded'].
    ↪min()
platform_range = data['platform_encoded'].max() - data['platform_encoded'].min()

random_genre_encoded = np.random.randint(data['genre_encoded'].min(),
    ↪data['genre_encoded'].min() + genre_range + 1)
random_publisher_encoded = np.random.randint(data['publisher_encoded'].min(),
    ↪data['publisher_encoded'].min() + publisher_range + 1)
random_platform_encoded = np.random.randint(data['platform_encoded'].min(),
    ↪data['platform_encoded'].min() + platform_range + 1)

new_random_data = {
    'genre_encoded': random_genre_encoded,
    'publisher_encoded': random_publisher_encoded,
    'platform_encoded': random_platform_encoded,
    'const': 1
}

new_random_data_df = pd.DataFrame([new_random_data])

actual_random_genre = data[data['genre_encoded'] ==
    ↪new_random_data['genre_encoded']]['genre'].iloc[0]
actual_random_publisher = data[data['publisher_encoded'] ==
    ↪new_random_data['publisher_encoded']]['publisher'].iloc[0]
actual_random_platform = data[data['platform_encoded'] ==
    ↪new_random_data['platform_encoded']]['platform'].iloc[0]

print("Actual Random Genre:", actual_random_genre)
print("Actual Random Publisher:", actual_random_publisher)
print("Actual Random Platform:", actual_random_platform)

predicted_random_global_sales = model.predict(new_random_data_df)

print("Predicted Random Global Sales:", predicted_random_global_sales.iloc[0])
```

Actual Random Genre: Shooter

Actual Random Publisher: Marvelous Entertainment

Actual Random Platform: PS2

Predicted Random Global Sales: 5.6818253263882585

Predictions made manually

```
[149]: new_data = {
        'genre_encoded': 10,
        'publisher_encoded': 359,
        'platform_encoded': 26,
        'const': 1
    }

    new_data_df = pd.DataFrame([new_data])

    # Get the original names for the encoded values
    actual_genre = data[data['genre_encoded'] ==
        ↪new_data['genre_encoded']]['genre'].iloc[0]
    actual_publisher = data[data['publisher_encoded'] ==
        ↪new_data['publisher_encoded']]['publisher'].iloc[0]
    actual_platform = data[data['platform_encoded'] ==
        ↪new_data['platform_encoded']]['platform'].iloc[0]

    # Print the actual categories
    print("Actual Genre:", actual_genre)
    print("Actual Publisher:", actual_publisher)
    print("Actual Platform:", actual_platform)

    # Make predictions
    predicted_global_sales = model.predict(new_data_df)

    print("Predicted Global Sales:", predicted_global_sales.iloc[0])
```

Actual Genre: Sports
 Actual Publisher: Nintendo
 Actual Platform: Wii
 Predicted Global Sales: 6.963548397872952

5.0.3 Statistical Analysis

Test for multicollinearity

- All VIF values are well below 5, which is a good sign. It suggests that there is no severe multicollinearity among the features.

```
[150]: vif_data = pd.DataFrame()
        vif_data["feature"] = X.columns

        vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                            for i in range(len(X.columns))]

        print(vif_data)
```

	feature	VIF
0	const	9.231023

```

1      genre_encoded  1.003293
2  publisher_encoded  1.002902
3  platform_encoded  1.000987

```

Test for heteroskedasticity

- Since the p-value (0.1400) is greater than the typical significance level (e.g., 0.05), we do not have sufficient evidence to reject the null hypothesis.
- This suggests that there is no strong evidence of heteroskedasticity in our model.

```

[151]: # Heteroskedasticity: Breush-Pagan --> Ho: var = constant

name = ["Lagrange multiplier statistic", "p-value", "f-value", "f p-value"]
test = sms.het_breuschpagan(model.resid, model.model.exog)
print(("BP Results:", ['bold']))
print(list(zip(name, test)))

('BP Results:', ['bold'])
[('Lagrange multiplier statistic', 5.476703548157174), ('p-value',
0.1400386619168625), ('f-value', 1.8257333842887522), ('f p-value',
0.14005412400102765)]

```

Evaluate the robustness of your coefficient estimates by bootstrapping model

```

[152]: ols_mod = sm.OLS(y, X).fit()

boot_slope_genre = []
boot_slope_publisher = []
boot_slope_platform = []
boot_interc = []
boot_adjR2 = []

n_boots = 100
n_points = data.shape[0]

plt.figure()

for _ in range(n_boots):
    sample_df = data.sample(n=n_points, replace=True)

    X_temp = sm.add_constant(sample_df[['genre_encoded', 'publisher_encoded',
↪ 'platform_encoded']])
    y_temp = sample_df['Global Sales']

    ols_model_temp = sm.OLS(y_temp, X_temp).fit()

    boot_interc.append(ols_model_temp.params[0])
    boot_slope_genre.append(ols_model_temp.params[1])
    boot_slope_publisher.append(ols_model_temp.params[2])

```

```

boot_slope_platform.append(ols_model_temp.params[3])
boot_adjR2.append(ols_model_temp.rsquared_adj)

y_pred_temp = ols_model_temp.predict(X_temp)
plt.plot(sample_df['Global Sales'], y_pred_temp, color='grey', alpha=0.2)

y_pred = ols_mod.predict(X)

plt.scatter(data['genre_encoded'], data['Global Sales'], label='Original Data:
↳genre_encoded')
plt.scatter(data['publisher_encoded'], data['Global Sales'], label='Original
↳Data: publisher_encoded')
plt.scatter(data['platform_encoded'], data['Global Sales'], label='Original
↳Data: platform_encoded')

plt.plot(data['Global Sales'], y_pred, linewidth=2, color='red',
↳label='Regression Line')
plt.grid(True)
plt.xlabel('Values')
plt.ylabel('Global Sales')
plt.title('Global Sales vs. Various Variables')
plt.legend()
plt.show()

sns.displot(boot_interc, alpha=0.25)
plt.axvline(x=ols_mod.params[0], color='red', linestyle='--')
plt.axvline(x=np.mean(boot_interc), color='orange', linestyle='--')
plt.title('Bootstrap Estimates: Intercept')
plt.show()

sns.displot(boot_slope_genre, alpha=0.25)
plt.axvline(x=ols_mod.params[1], color='red', linestyle='--')
plt.axvline(x=np.mean(boot_slope_genre), color='orange', linestyle='--')
plt.title('Bootstrap Estimates: Slope for genre_encoded')
plt.show()

sns.displot(boot_slope_publisher, alpha=0.25)
plt.axvline(x=ols_mod.params[2], color='red', linestyle='--')
plt.axvline(x=np.mean(boot_slope_publisher), color='orange', linestyle='--')
plt.title('Bootstrap Estimates: Slope for publisher_encoded')
plt.show()

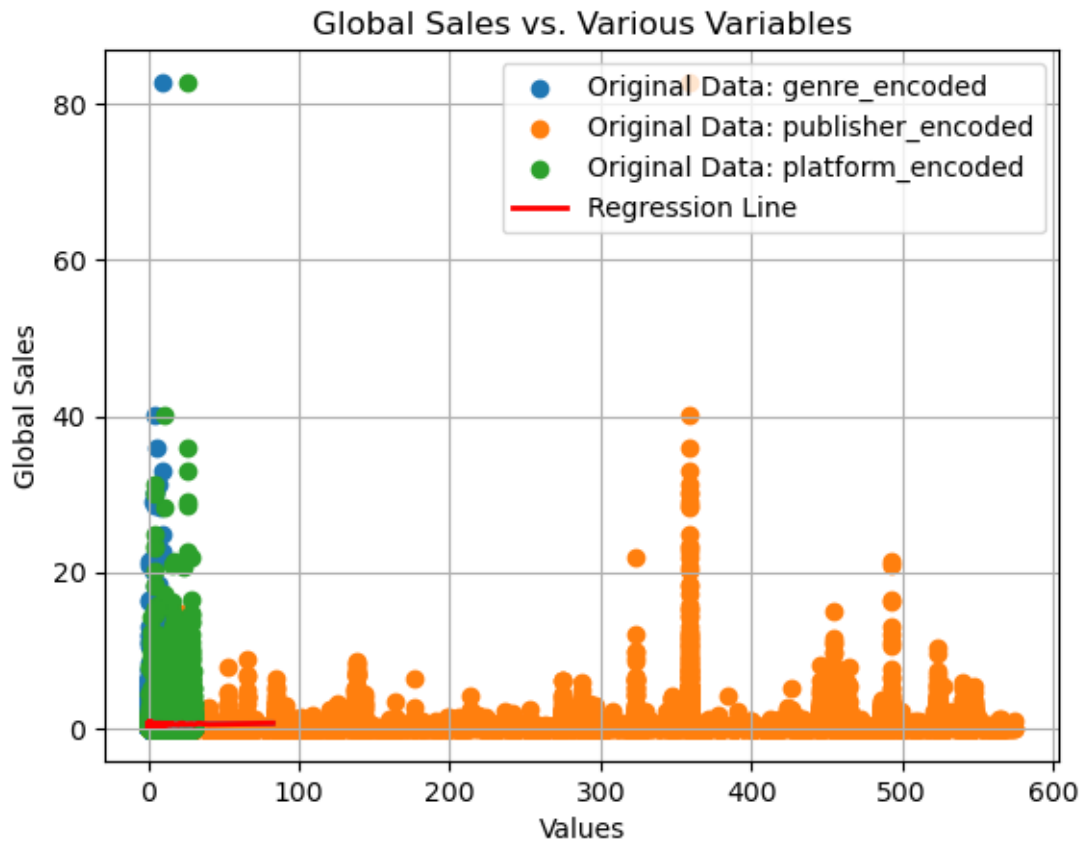
sns.displot(boot_slope_platform, alpha=0.25)
plt.axvline(x=ols_mod.params[3], color='red', linestyle='--')
plt.axvline(x=np.mean(boot_slope_platform), color='orange', linestyle='--')
plt.title('Bootstrap Estimates: Slope for platform_encoded')
plt.show()

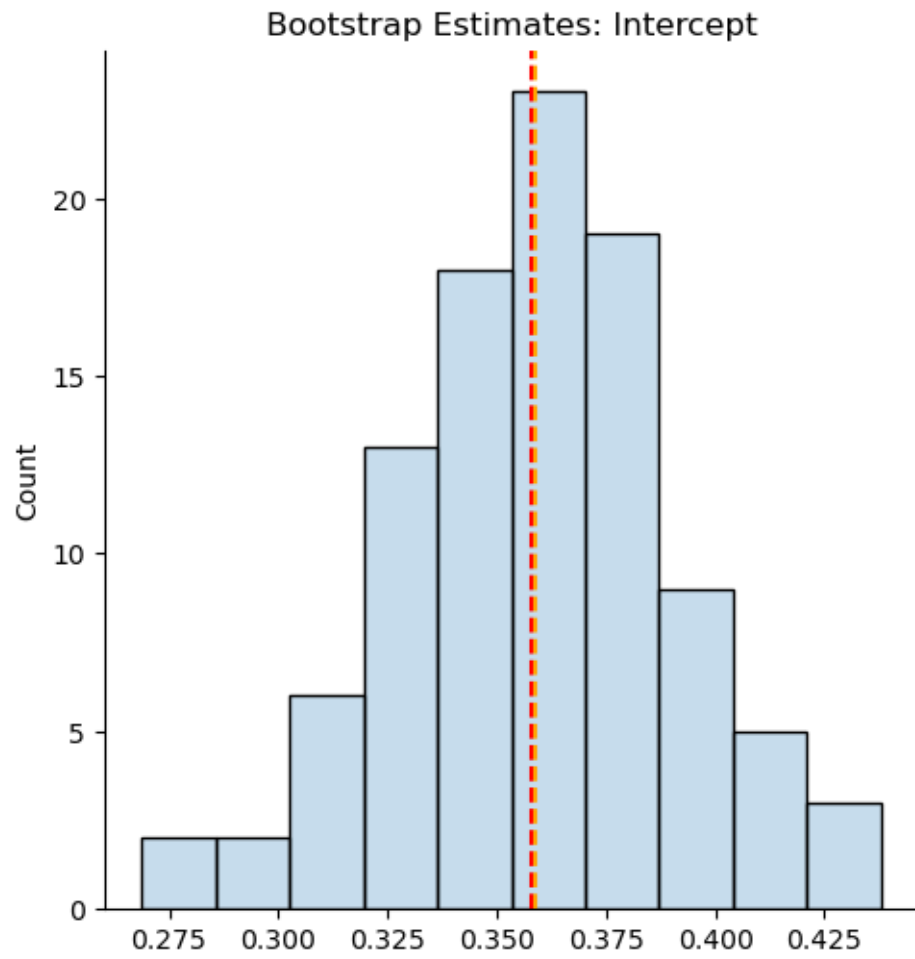
```

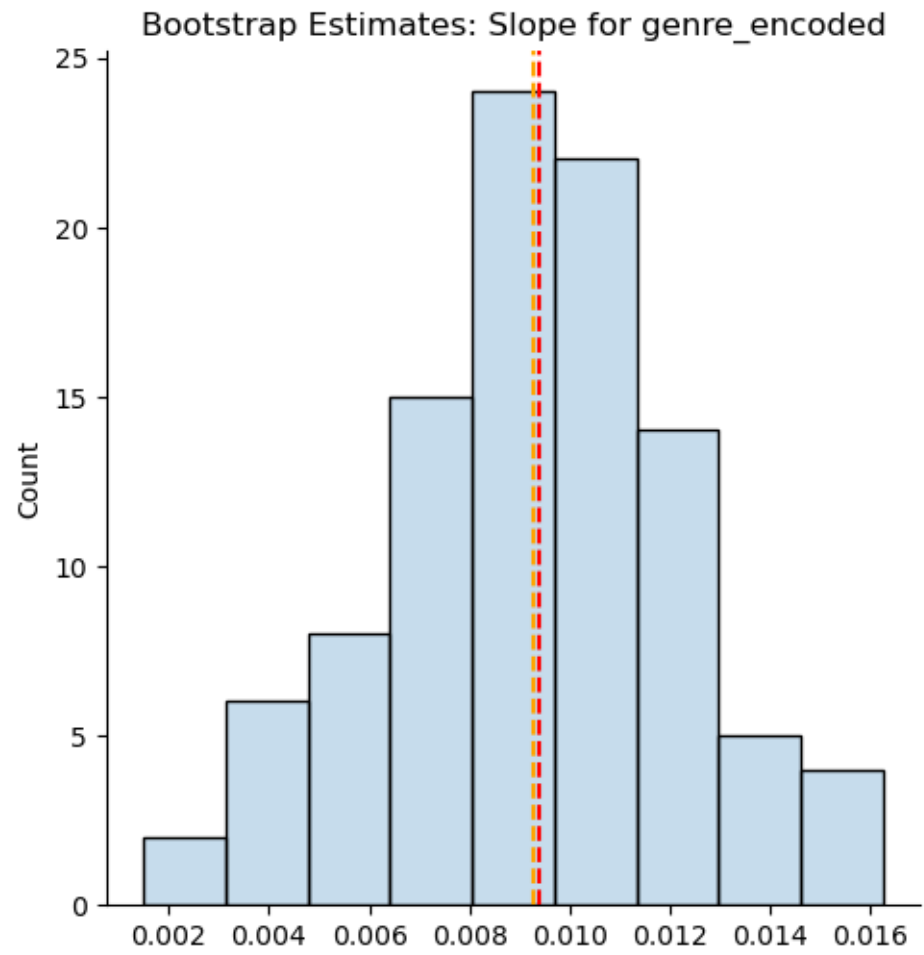
```

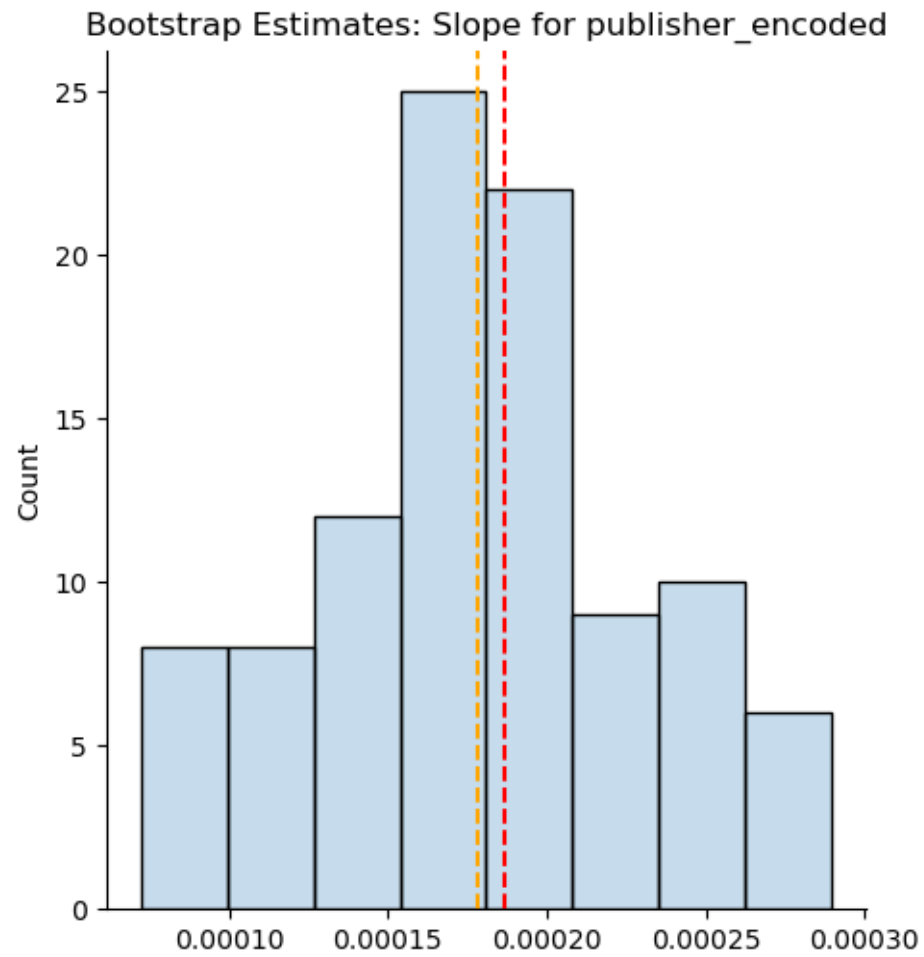
sns.displot(boot_adjR2, alpha=0.25)
plt.axvline(x=ols_mod.rsquared_adj, color='red', linestyle='--')
plt.axvline(x=np.mean(boot_adjR2), color='orange', linestyle='--')
plt.title('Bootstrap Estimates: Adjusted R-squared')
plt.show()

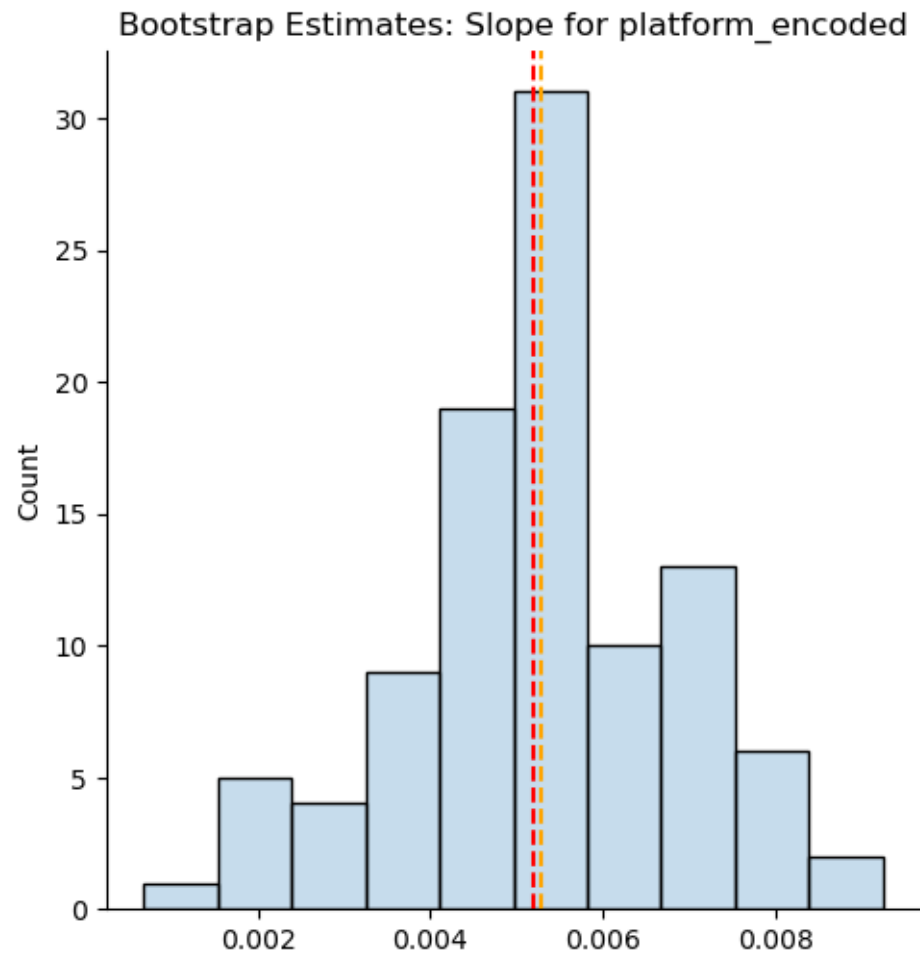
```

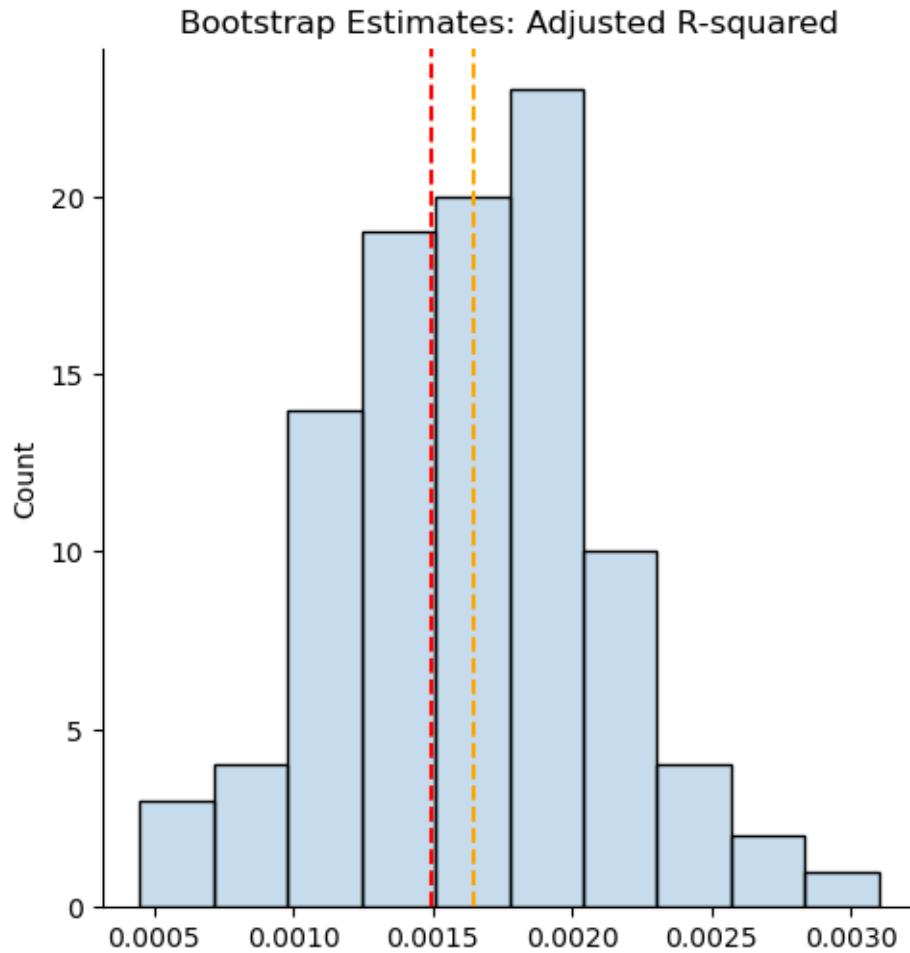












Use cross-validation to evaluate your model's performance

```
[153]: regr = LinearRegression()
model_val = regr.fit(X,y)
regr.coef_
regr.intercept_

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=0)

regr = LinearRegression()
regr.fit(x_train, y_train)

y_pred = regr.predict(x_test)
```

```
print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

regr = linear_model.LinearRegression()
scores = cross_val_score(regr, X, y, cv=5,
    ↳scoring='neg_root_mean_squared_error')
print('5-Fold CV RMSE Scores:', scores)
```

MAE: 0.5924579386055017

MSE: 1.9563671461045857

RMSE: 1.398701950418525

5-Fold CV RMSE Scores: [-3.58691194 -0.22708645 -0.47090875 -0.57687974
-0.63498612]

[]: